

Data Engineering for Data Solutions
Case study
Candidate version

Business problem

- Parcel service (delivery)
- Increasing number of complaints
- Root cause is nuclear
- Tight monitoring to get hotspots

Deliverables

- Report 1:
 - What are the pick-up stations that received complaints in the last 14 days?
 - Pickup Station ID
 - Pickup Station Name
 - Pickup Station Zip code
 - Number of claims
- Report 2:
 - What are the final recipients and the associated stations with more than 5 complaints in last 2 months?
 - Final recipient name
 - Final recipient address
 - Pickup Station IDs

Goals

- Design a data pipeline
- Bring the delivery events data into the pipeline
- Bring the complaints messages into the pipeline
- Use a Cloud Service to aid in the process

Available data

- Delivery events files (CSV, XML and JSON)
 - Delivery timestamp
 - Pickup station ID
 - Final recipient's name and address
 - Driver ID
 - Parcel ID
 - Purchase Order ID
- Complaints (CSV, XML and JSON)
 - Timestamp of the complaint
 - Pickup station ID
 - Final recipient's name and address
 - Parcel ID (optional)
 - Purchase Order ID (optional)
- List of Pickup stations (database)
 - Pickup station ID
 - Name
 - City
 - Zip code
 - Street

Assumptions

- We are using Microsoft Azure as Cloud Service
- Azure PostgreSQL as database management system
- Azure Data Factory as orchestration tool and Azure Data Lake Storage Gen2 as data lake
- Pickup IDs in complaints data are congruent with the reference database
- Pickup IDs in delivery events data are not consolidated with reference database
- Delivery events and complaints files fall into two different folders in the central server
- Pipeline will run once a day

Pipeline architecture

- Three layers for improved control and governance
 - Landing zone
 - Staging zone
 - Data Warehouse zone
- Orchestration will be made mostly by pipelines in Azure Data Factory
- Reports will be exposed as views

1) bring raw data into landing zone

- Azure Data Factory Pipelines
 - Daily triggered
 - Get metadata
 - ForEach file
 - Copy activity
 - SQL Script (INSERT)
- Copies raw data into Azure Data Lake Storage Gen2
- Updates *landing_file* control table

2) Parse raw CSV/XML/JSON into staging zone

- Azure Data Factory Pipeline
 - Trigger
 - LookUp “RECEIVED” in *landing_file*
 - If COUNT greater than zero
 - Calls Python parser function or container
- Python parser inserts data into suitable staging table (*stg_delivery_events* or *stg_complaints*)
- For each parsed file updates *landing_file* with “PARSED” or “FAILED”

3) copy reference data into staging zone

- Azure Data Factory Pipelines
 - Trigger
 - SQL Script (INSERT)
- Creates a fingerprint to be used in later JOINs
- Creates a HASH with name, city, zip code and street to easily test if anything has changed

4) update SCD2 *dim_pickup_station* in dw zone

- Azure Data Factory Pipeline
 - Trigger
 - SQL Script (ADD)
 - SQL Script (UPDATE CLOSE)
 - SQL Script (UPDATE INSERT)
 - SQL Script (REMOVE)
- Update the SCD2 by adding new entries, closing and inserting entries that changed and deactivating removed entries.
- Creates a canonical ID.

5) Update fact tables in dw zone

- Azure Data Factory Pipeline
 - Trigger
 - SQL Script (INSERT)
- Inserts data from *staging.stg_delivery_events* into *dw.fact_delivery_events*
- Inserts data from *staging.stg_complaints* into *dw.fact_complaints* joining with *dw.dim_pickup_station* via Pickup Station ID

6) monitoring data integration

- Azure Data Factory Pipeline
 - Trigger
 - SQL Script (INSERT)
- Inserts into *dw.monitoring_complaints* the number os complaints processed in the current day and the number of complaints that didn't get a pickup ID match with *dw.dim_pickup_station*.

7) reports

- Once they are views, they are instantly updated
- Report 1 is exposed in *dw. rpt_pickup_station_claims_14d*
- Report 2 is exposed in *dw. rpt_recipients_over5_complaints_2m*

Next steps

- Improve monitoring stages in pipeline
- Correlate complaints and delivery events using *parcel_id* or *purchase_order_id* whenever possible
- Create a crosswalk table allowing a more precise connection between delivery events and complaints
- Draw insights from drivers, recipients and stations in terms of number of deliveries and number of complaints
- Use LLM to evaluate the claims contents